

Towards a Better Verification function

Lixun Zhang

August 17, 2022

Outline

1 Motivation

2 Experiment

- Methodology and Setup
- Results and Observations
- Summary

3 Proposals

Verification Method in MLIR

Verification: compare outputs from GPU kernel (kern) and validation function (val)

Current verification method used by MLIR

$$\text{maxRelDiff}_{\text{old}} = \max_i \left(\frac{|\text{val}_i - \text{kern}_i|}{|\text{val}_i|} \right), |\text{val}_i| > 1 \times 10^{-3}$$

- Element wise metric
- Relative error between two values
- Ignore the error if the denominator is too small ($< 1 \times 10^{-3}$)
- Pass if the maximum relative difference among all elements is smaller than a tolerance δ_{rel}

Problem: $\delta_{rel} = 15\%$ for fp16 tests with random inputs $([-1, 1])$

- MLIR#408, PR#694, PR#696
- δ_{rel} is too large to distinguish bugs from numerical errors
- For other data types, $\delta_{rel} = 1 \times 10^{-6}$

Verification Method in MIOpen

Current verification method used by MIOpen: RMS

- global metric
- Pass if RMS is smaller than a tolerance δ_{RMS}
- Adjust δ_{RMS} for different directions, algorithms, data types, and issues

$$\frac{\sqrt{\sum_{i=1}^N (a_i - b_i)^2}}{\sqrt{N} \cdot \max(\max(|a_i|), \max(|b_i|))}$$

| | | int32 | int8 | fp32 | fp16 | bf16 |
|------------|----------|----------|----------|------------------------------|----------|----------|
| fwd bwd | other | 1.50E-06 | 1.50E-06 | 1.50E-06 | 8.20E-03 | 6.56E-02 |
| | igemm | 1.50E-05 | 1.50E-05 | 1.50E-05 | 8.20E-02 | 0.656 |
| wrw | other | 3.00E-06 | 3.00E-06 | 3.00E-06 | 1.64E-02 | 0.13 |
| | winograd | 3.00E-06 | 3.00E-06 | 3.00E-05 | 1.64E-02 | 0.13 |
| | igemm | 3.00E-06 | 3.00E-06 | 3.00E-05 0.01 (issue2176) | 8.20E-02 | 0.13 |

Pro/con of RMS as a measure of difference

- Insensitive to element wise differences

Element wise vs. RMS

Element wise metric

- compare each element
- pick the maximum difference
- sensitive to large differences for individual elements
- false alarm caused by numerical errors

Aggregated global metric

- compare each element
- aggregate differences from all elements
- insensitive to large differences for individual elements
- false negative and miss a real bug

We need a better verification function that can **detect bugs and ignore numerical errors**

- which metric to use? Element wise or RMS?
- how to choose δ ?
- how to deal with special tests?

Other Element Wise Metrics

- Absolute difference $\text{maxAbsDiff} = \max_i (|\text{val}_i - \text{kern}_i|)$
- Relative difference (ignore zero denominators)

$$\text{maxRelDiff} = \max_i \left(\frac{|\text{val}_i - \text{kern}_i|}{|\text{val}_i|} \right), |\text{val}_i| > 0$$

- Relative difference w.r.t ϵ

$$\text{maxEpsilonDiff} = \max_i \left(\frac{|\text{val}_i - \text{kern}_i|}{\epsilon_{|\text{val}_i|}} \right)$$

where $\epsilon_{|\text{val}_i|}$ is the precision of fp16 numbers around val_i .

Experiment Setup

Sources of configs: auto_e2e/ e2e_for_pr/ misc_e2e/

| | fixed | random | PR CI | Nightly CI |
|------------|-------------|--------|-------|------------|
| e2e_for_pr | pv_with_gpu | | | |
| auto_e2e | pv_with_gpu | pv | | |
| misc_e2e | pv_with_gpu | | | |

Hardware

- MI100 gfx908 ixt-rack-112
- MI200 gfx90a lockhart5

Settings of the test

- -rand 1: Generate random numbers as inputs
- -rand_min and -rand_max: choose the range of the random numbers
 - r0 [-1, 1], r1 [-10, 10], r4 [1, 5], r5 [5, 10]
- -pv and -pv_with_gpu: choose the validation method
- -x2: enable xdlops

Methodology: Effects of different settings to the error w.r.t all metrics

Flush-denorms-to-zero

| xdlops/cpu | MI200 | | | MI100 | | |
|--------------------|----------|----------|----------|----------|----------|----------|
| | r0 | r1 | r4 | r0 | r1 | r4 |
| maxEpsilon ave | 423.43 | 986.77 | 0.87 | 74.89 | 1,084.96 | 0.86 |
| maxEpsilon max | 2446 | 13389 | 1 | 2012 | 18788 | 1 |
| maxAbsDiff ave | 0.04 | 3.31 | 3.08 | 0.04 | 3.16 | 3.08 |
| maxAbsDiff max | 0.5 | 32 | 32 | 0.5 | 32 | 32 |
| maxRelDiff ave | 6.33 | 1.71 | 6.85E-04 | 1.63 | 1.53 | 6.82E-04 |
| maxRelDiff max | 230 | 44 | 9.80E-04 | 20 | 25 | 9.80E-04 |
| maxRelDiff old ave | 0.03 | 0.07 | 6.85E-04 | 4.30E-03 | 0.07 | 6.82E-04 |
| maxRelDiff old max | 0.077 | 1.2 | 9.80E-04 | 0.051 | 0.97 | 9.80E-04 |
| RMS ave | 2.08E-06 | 1.12E-06 | 4.79E-06 | 1.41E-06 | 1.11E-06 | 4.75E-06 |
| RMS max | 1.40E-05 | 3.80E-06 | 2.30E-05 | 1.40E-05 | 4.40E-06 | 2.30E-05 |

mfma instructions flush subnormal numbers ($< 2^{-14}$) to zero on MI200
 r0 [-1, 1]
 r1 [-10, 10]
 r4 [1, 5]

- flush-denorms-to-zero behavior leads to
 - $E_{r1} > E_{r0}$ on both MI200 and MI100 (except for maxRelDiff and RMS)
 - $E_{MI200} > E_{MI100}$ with r0
 - $E_{MI200} \approx E_{MI100}$ with r4

Underflow and Overflow

| nonxdlops/ cpu | MI200 | | |
|--------------------|----------|----------|----------|
| | r0 | r4 | r5 |
| maxEpsilon ave | 48.13 | 0.36 | 0.42 |
| maxEpsilon max | 1367 | 1 | 1 |
| pass rate | 68.22% | 100.00% | 99.07% |
| maxAbsDiff ave | 0.012 | 1.322 | 2.755 |
| maxAbsDiff max | 0.5 | 32 | 32 |
| maxRelDiff ave | 0.75 | 2.84E-04 | 2.57E-04 |
| maxRelDiff max | 18 | 9.80E-04 | 9.80E-04 |
| maxRelDiff old ave | 2.31E-03 | 2.84E-04 | 2.57E-04 |
| maxRelDiff old max | 3.80E-02 | 9.80E-04 | 9.80E-04 |
| RMS ave | 6.20E-07 | 3.39E-06 | 2.08E-06 |
| RMS max | 3.50E-06 | 2.60E-05 | 1.40E-05 |


r0 [-1, 1] vs. r4 [1, 5]/r5 [5, 10]

A test passes if $\text{maxEpsilonDiff} \leq 1$

- Errors with r0 is larger than that of r4/r5
 \Leftarrow Underflow
- Except for maxAbsDiff and RMS
 \Leftarrow larger inputs with r4/r5
- All tests with r4 pass since the maximum maxEpsilonDiff is 1
- However, some tests with r5 fail even the maximum maxEpsilonDiff is 1
 \Leftarrow Overflow (> 65504)

Validation Methods

| kernel/ validation | MI200 | | | | | |
|-----------------------|-------------------|----------------|----------------------|-------------------|----------------|----------------------|
| | r0 | | | r4 | | |
| | nonxdlops/ cpu | xdlops/ cpu | xdlops/ nonxdlops | nonxdlops/ cpu | xdlops/ cpu | xdlops/ nonxdlops |
| maxEpsilon ave | 48.13 | 423.43 | 471.15 | 0.36 | 0.87 | 0.95 |
| maxEpsilon max | 1367 | 2446 | 2278 | 1 | 1 | 1 |
| maxAbsDiff ave | 0.012 | 0.04 | 0.405 | 1.32 | 3.08 | 3.61 |
| maxAbsDiff max | 0.5 | 0.5 | 16 | 32 | 32 | 32 |
| maxRelDiff ave | 0.75 | 6.33 | 6.17 | 2.84E-04 | 6.85E-04 | 7.24E-04 |
| maxRelDiff max | 18 | 230 | 86 | 9.80E-04 | 9.80E-04 | 9.80E-04 |
| maxRelDiff old ave | 2.31E-03 | 2.63E-02 | 3.33E-02 | 2.84E-04 | 6.85E-04 | 7.24E-04 |
| maxRelDiff old max | 0.038 | 0.077 | 0.17 | 9.80E-04 | 9.80E-04 | 9.80E-04 |
| RMS ave | 6.20E-07 | 2.08E-06 | 3.67E-06 | 3.39E-06 | 4.79E-06 | 7.84E-06 |
| RMS max | 3.50E-06 | 1.40E-05 | 2.90E-05 | 2.60E-05 | 2.30E-05 | 2.50E-05 |



r0 [-1, 1] vs. r4 [1, 5]

- cpu: sequential implementation of the convolution operation
- non-xdlops: convert convolution to gemm and compute gemm without `mfma`
- xdlops: convert convolution to gemm and compute gemm with `mfma`

Exploration of a “Better” Validation Function

MI200 xdlops/cpu

auto_e2e/padding_kernel_gemmK.mlir [CHECK_RESNET50_F16_CONFIG1]

| relDiff old | original | acc 4 | flush+acc 4 |
|--------------|----------------|----------------|----------------|
| 0 | 801676(99.86%) | 802015(99.90%) | 802369(99.94%) |
| (0, 1e-6) | 0(0.0%) | 0(0.0%) | 0(0.0%) |
| [1e-6, 1e-5) | 0(0.0%) | 0(0.0%) | 0(0.0%) |
| [1e-5, 1e-4) | 0(0.0%) | 0(0.0%) | 0(0.0%) |
| [1e-4, 1e-3) | 970(0.12%) | 704(0.09%) | 364(0.04%) |
| [1e-3, 1e-2) | 60(0.007474%) | 60(0.007474%) | 0(0.0%) |
| [1e-2, 0.1) | 3(0.000374%) | 3(0.000374%) | 0(0.0%) |
| [0.1, 1) | 0(0.0%) | 0(0.0%) | 0(0.0%) |
| >= 1 | 0(0.0%) | 0(0.0%) | 0(0.0%) |
| val == 0 | 0(0.0%) | 0(0.0%) | 0(0.0%) |
| max | 0.027 | 0.028 | 0.00097 |
| val | -0.001073837 | -0.001074791 | 0.125244141 |
| kern | -0.001045227 | -0.001045227 | 0.125366211 |

Push CPU to emulate the behavior of the xdlops pipeline

- acc 4: the mfma instructions are performed by the DOT4_F32_F16 unit, which performs a fused multiply and add of 4 pairs of fp16 numbers
⇒ slightly improves the results
- flush: mfma instructions on MI200
⇒ greatly improves the results

acc 4: truncate the sum of every 4 multiplications from fp64 to fp32

flush: flush small numbers ($< 2^{-14}$) to zero

Summary

Causes of numerical errors for fp16 tests:

- ① Subnormal numbers flushed to zero
- ② Computation order difference due to partition of gemm onto the grid
- ③ Truncation behavior difference during accumulation of intermediate results

Element wise vs. RMS

- Element wise metrics are more sensitive to numerical errors
⇒ not suitable for small inputs
- RMS is more sensitive to the magnitude of the elements
⇒ not suitable for large inputs

Upcoming Testing Framework

- ❶ Combine RMS and element wise metrics together (MLIR#620)
 - For random inputs with $[-1, 1]$, use RMS
 - For random inputs with $[-3, -1]$ and $[1, 3]$, use maxRelDiff
- ❷ Let each test decide how to verify the results (MLIR#621)
 - choose the range of random numbers
 - choose which metric to use
 - choose the tolerance for each metric
 - RMS: 3×10^{-5}
 - maxRelDiff: 1×10^{-3}